Supplementary Material

Heechul Yang, Jaewoo Kang*

Korea University
kangj@korea.ac.kr

Models and policies

The main pool has small ($\leq$14B parameters) and mid-sized models; large models are reported separately and are not part of the main mean or the improvement criterion. All models run in Instant mode (no thinking) with greedy decoding. Policies: C0 (answer only) answers immediately; the free agent acts for up to five turns among ANSWER, ASK, SEARCH and ABSTAIN with at most two questions and two searches; forced inquiry must ask or search in each of the first four turns and answer or abstain at the fifth; the targeted ask pipeline writes three questions at once, instructed to ask for the patient facts that separate the two most plausible options, and answers once with all replies. The noticing probe appends "Is the information in this case sufficient to answer the question? Answer Yes or No." to each intact and fact-removed case and scores it by the next-token logit of No minus that of Yes; the AUROC treats fact-removed cases as positives and intact cases as negatives (rank-based, ties counted half). All intervals are 95% item-clustered paired bootstrap intervals ($B = 10{,}000$). Every value in these tables is read from results/p2_analysis.json and the analysis files next to it.

Results by model

This section gives, for every model, the cost of removal (Figure S1), accuracy in each policy and state, the improvement criterion with its sensitivity analyses, free-agent and forced-inquiry behavior, acquisition, the oracle and the noticing probe, the fine-tune versus base comparisons, and output format.

Additional analyses

The analyses in this section are exploratory or post hoc, as marked in each caption; no multiplicity correction is applied.

Benchmark and simulator checks

The fact-type labels are properties of the items. The refusal audit sampled questions logged in runs of an earlier model set with the same simulator, not questions from the models evaluated here; it describes the simulator, not those models. Table S12 gives the construction counts and necessity labels.

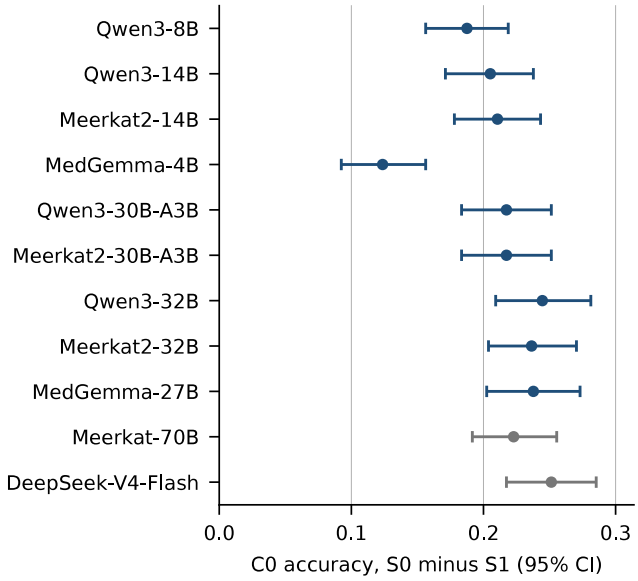


Figure S1: Cost of removal: C0 (answer only) accuracy drop from intact (S0) to fact-removed (S1) cases, with 95% CIs. Grey: large models, reported separately.

*Corresponding author.

Table S1: Accuracy by model, policy and state (736 items). C0 has no simulator, so its S2 input equals S1. Large models are reported separately.

	C0		Free agent		Forced inquiry		Targeted ask	
Model	S0	S1	S1	S2	S1	S2	S1	S2
Qwen3-8B	0.598	0.410	0.452	0.424	0.438	0.413	0.432	0.390
Qwen3-14B	0.637	0.432	0.424	0.413	0.471	0.425	0.470	0.420
Meerkat2-14B	0.735	0.524	0.515	0.515	0.505	0.489	0.558	0.515
MedGemma-4B	0.515	0.391	0.394	0.391	0.399	0.391	0.390	0.387
Qwen3-30B-A3B	0.704	0.486	0.477	0.477	0.496	0.467	0.523	0.478
Meerkat2-30B-A3B	0.765	0.548	0.539	0.538	0.569	0.550	0.576	0.537
Qwen3-32B	0.717	0.473	0.486	0.481	0.523	0.481	0.539	0.467
Meerkat2-32B	0.755	0.519	0.507	0.507	0.527	0.519	0.590	0.520
MedGemma-27B	0.719	0.481	0.500	0.488	0.579	0.516	0.526	0.474
Meerkat-70B	0.720	0.497	0.492	0.492	0.562	0.478	0.550	0.495
DeepSeek-V4-Flash	0.804	0.553	0.571	0.568	0.632	0.622	0.663	0.545

95% CIs for every cell are in results/p2_analysis.json.

Table S2: Improvement criterion fixed before the runs: mean Δ at S1 at least $+0.05$, lower bound of the 97.5% interval above zero, and mean change at S2 at least -0.02, on the main mean of the final pool (first row of each policy). The other rows are sensitivity analyses; earlier pool: the main pool before Meerkat2-32B was added, shown for reference only.

Policy	Pool	Δ at S1 [95% CI]	97.5% CI	Δ at S2 [95% CI]	Criterion
Forced inquiry	main, 9 models	$+0.027$ [$+0.016$, $+0.038$]	[$+0.014$, $+0.040$]	-0.001 [-0.011, $+0.009$]	not met
	without MedGemma-4B, 8	$+0.029$ [$+0.018$, $+0.041$]	[$+0.016$, $+0.043$]	-0.002 [-0.012, $+0.009$]	not met
	earlier pool, 8	$+0.029$ [$+0.017$, $+0.041$]	[$+0.016$, $+0.043$]	-0.002 [-0.012, $+0.009$]	not met
Targeted ask	main, 9 models	$+0.038$ [$+0.029$, $+0.047$]	[$+0.027$, $+0.049$]	-0.008 [-0.015, -0.002]	not met
	without MedGemma-4B, 8	$+0.043$ [$+0.033$, $+0.053$]	[$+0.032$, $+0.055$]	-0.009 [-0.016, -0.002]	not met
	valid sets only (632 items)	$+0.039$ [$+0.029$, $+0.050$]	[$+0.027$, $+0.051$]	-0.010 [-0.017, -0.003]	not met
	earlier pool, 8	$+0.034$ [$+0.024$, $+0.043$]	[$+0.023$, $+0.044$]	-0.010 [-0.016, -0.003]	not met

Table S3: Behavior (% of episodes). The first action is taken before any simulator reply. Abstained episodes are completed with a forced answer for scoring.

	Free agent, first action at S1				Free agent abstains		Forced inquiry abstains	
Model	ANSWER	ASK	SEARCH	ABSTAIN	S1	S2	S1	S2
Qwen3-8B	6.7	1.9	91.2	0.3	2.0	2.3	1.1	1.1
Qwen3-14B	60.7	0.7	38.3	0.3	8.6	9.5	21.6	25.8
Meerkat2-14B	90.1	0.1	9.1	0.7	3.8	3.8	45.8	52.0
MedGemma-4B	68.2	18.8	12.9	0.0	0.1	0.1	7.1	9.5
Qwen3-30B-A3B	58.3	0.1	41.6	0.0	2.7	2.7	10.9	14.7
Meerkat2-30B-A3B	32.7	2.6	64.4	0.0	4.5	4.6	45.9	52.0
Qwen3-32B	14.1	1.5	83.7	0.7	9.2	9.6	15.4	22.8
Meerkat2-32B	87.0	0.0	11.8	1.1	3.5	3.5	33.7	35.5
MedGemma-27B	34.9	1.0	64.1	0.0	0.5	1.4	0.7	1.6
Meerkat-70B	99.3	0.0	0.7	0.0	0.0	0.0	0.0	0.3
DeepSeek-V4-Flash	97.7	1.0	1.4	0.0	0.0	0.0	0.0	0.0

Table S4: Acquisition at S1. Q/item: questions asked per item. Disclosed: % of questions to which the simulator disclosed a removed fact. $\geq$1: % of items that received at least one removed fact. –: no question asked.

	Free agent			Forced inquiry			Targeted ask		
Model	Q/item	disclosed	$\geq$1	Q/item	disclosed	$\geq$1	Q/item	disclosed	$\geq$1
Qwen3-8B	0.45	28.4	12.5	0.58	26.5	15.5	3.00	18.2	43.8
Qwen3-14B	0.23	11.9	2.6	1.49	19.2	24.9	3.00	20.4	48.2
Meerkat2-14B	0.01	25.0	0.3	0.57	18.9	9.9	2.84	25.4	54.1
MedGemma-4B	0.30	12.4	3.5	1.88	9.2	11.3	2.96	11.2	28.1
Qwen3-30B-A3B	0.00	33.3	0.1	0.52	25.5	13.0	2.99	22.5	51.0
Meerkat2-30B-A3B	0.03	36.4	1.1	1.35	11.2	12.8	2.76	24.6	48.4
Qwen3-32B	0.07	26.9	1.9	0.74	33.6	21.6	2.96	26.3	59.2
Meerkat2-32B	0.00	100.0	0.1	0.15	41.7	5.3	2.50	30.7	55.4
MedGemma-27B	0.16	42.2	6.2	0.99	35.8	34.0	3.00	23.3	55.2
Meerkat-70B	0.00	–	0.0	1.85	24.0	35.7	2.71	24.4	50.1
DeepSeek-V4-Flash	0.01	66.7	0.8	0.05	57.9	3.0	3.00	29.0	66.3

Table S5: Handing over the removed facts and the information-sufficiency probe. R = (accuracy with both facts − C0 S1) / (C0 S0 − C0 S1); reading fixed before the run: at least 0.80, the facts restore most of the lost accuracy. One fact: accuracy with only the first or the second removed fact. NP AUROC: fact-removed (S1) against intact (S0) versions of the same items. No equivalence test was run. Main-mean NP AUROC by necessity label (post hoc, uncorrected): 0.72 [0.71, 0.74] necessary, 0.70 [0.69, 0.73] partially necessary, 0.67 [0.64, 0.70] not necessary, 0.70 [0.69, 0.72] all items; differences between strata were not formally tested.

Model	R [95% CI]	Both facts − S0 [95% CI]	One fact (F1)	One fact (F2)	NP AUROC [95% CI]
Qwen3-8B	1.06 [0.94, 1.19]	+0.011 [−0.012, +0.033]	0.554	0.531	0.67 [0.66, 0.69]
Qwen3-14B	1.03 [0.93, 1.14]	+0.005 [−0.016, +0.026]	0.561	0.554	0.71 [0.70, 0.73]
Meerkat2-14B	0.98 [0.89, 1.08]	−0.004 [−0.023, +0.016]	0.674	0.637	0.74 [0.72, 0.76]
MedGemma-4B	1.14 [0.96, 1.38]	+0.018 [−0.005, +0.041]	0.495	0.476	0.65 [0.64, 0.67]
Qwen3-30B-A3B	1.01 [0.92, 1.12]	+0.003 [−0.018, +0.024]	0.647	0.602	0.67 [0.66, 0.69]
Meerkat2-30B-A3B	1.02 [0.93, 1.12]	+0.004 [−0.015, +0.023]	0.705	0.668	0.73 [0.71, 0.74]
Qwen3-32B	0.97 [0.88, 1.08]	−0.007 [−0.031, +0.018]	0.633	0.625	0.73 [0.72, 0.75]
Meerkat2-32B	1.03 [0.95, 1.12]	+0.007 [−0.012, +0.026]	0.705	0.677	0.77 [0.75, 0.78]
MedGemma-27B	0.83 [0.74, 0.93]	−0.039 [−0.064, −0.015]	0.615	0.588	0.66 [0.65, 0.68]
Meerkat-70B	1.03 [0.94, 1.13]	+0.007 [−0.014, +0.029]	0.659	0.640	0.66 [0.65, 0.68]
DeepSeek-V4-Flash	0.99 [0.91, 1.08]	−0.001 [−0.023, +0.020]	0.739	0.717	0.64 [0.62, 0.66]

Table S6: Medical fine-tunes against their Qwen3 base models (paired over items, 95% item-clustered bootstrap). Columns: size of the pair (Meerkat2 fine-tune minus its Qwen3 base). The comparisons were specified in the plan without a decision rule; the accuracy and cost rows are descriptive (no equivalence test).

Meerkat2 − Qwen3	14B	30B-A3B	32B
C0 accuracy at S0	+0.098 [+0.068, +0.128]	+0.061 [+0.034, +0.090]	+0.038 [+0.011, +0.065]
Cost of removal (C0, S0 − S1)	+0.005 [−0.030, +0.042]	+0.000 [−0.035, +0.035]	−0.008 [−0.043, +0.027]
Δ at S1, free agent	−0.001 [−0.030, +0.027]	+0.001 [−0.026, +0.030]	−0.026 [−0.053, +0.001]
Δ at S1, forced inquiry	−0.058 [−0.098, −0.019]	+0.012 [−0.023, +0.048]	−0.042 [−0.077, −0.007]
Δ at S1, targeted ask	−0.004 [−0.034, +0.024]	−0.008 [−0.034, +0.018]	+0.004 [−0.027, +0.037]
Free agent opens with ANSWER	+0.293 [+0.258, +0.329]	−0.255 [−0.298, −0.213]	+0.728 [+0.696, +0.760]
Forced inquiry abstains	+0.242 [+0.200, +0.284]	+0.351 [+0.311, +0.390]	+0.183 [+0.145, +0.223]
Disclosed per question, targeted ask	+0.051 [+0.031, +0.071]	+0.021 [−0.001, +0.043]	+0.044 [+0.022, +0.065]

Table S7: Output format, summed over the C0, free-agent and forced-inquiry cells at S0, S1 and S2. Generation stops at the end of the first JSON object, a rule fixed before the runs. Unparseable: final answers that could not be parsed; Rescued: of those, answers that start with an option letter followed by the option text, accepted for every model; Letter+text: all outputs of that form. Invalid sets: malformed three-question sets of the targeted ask pipeline (including re-prompts).

Model	Outputs	$\geq$2 JSON	Truncated	Unparseable	Rescued	Letter+text	Invalid sets
Qwen3-8B	15,145	0	0	1	0	1,034	0 / 1,472
Qwen3-14B	12,509	0	0	0	0	255	0 / 1,472
Meerkat2-14B	11,272	0	0	0	0	37	8 / 2,606
MedGemma-4B	11,950	0	0	3	0	0	18 / 1,492
Qwen3-30B-A3B	11,954	0	0	0	0	7	4 / 1,490
Meerkat2-30B-A3B	12,146	0	0	10	0	23	44 / 2,630
Qwen3-32B	12,560	0	0	0	0	0	0 / 1,472
Meerkat2-32B	11,311	0	0	9	0	0	140 / 2,730
MedGemma-27B	13,478	0	4	2	0	110	0 / 2,004
Meerkat-70B	15,422	0	2	2101	2098	6,446	144 / 2,920
DeepSeek-V4-Flash	11,109	0	0	0	0	0	0 / 1,472

MedGemma-4B outputs come from an earlier run of the same frozen pipeline (same items, simulator, prompts and scoring), and MedGemma 1.5 4B was dropped after failing the format pilot: under C0 it wrote explanations instead of an option letter in the answer field, so even after a re-prompt 78% of S0 cases had no usable answer, and its targeted-ask question sets began with a thought marker and were cut off at 200 tokens, leaving 50/50 invalid. For MedGemma-27B the thought-marker tokens were blocked during generation. DeepSeek-V4-Flash's forced-inquiry runs were split across two parallel instances; re-running 10 multi-turn items in both states on one instance reproduced the stored outputs exactly.

Table S8: Cost of removal and Δ at S1 by necessity stratum (LLM labels): N necessary (380 items), P partially necessary (251), NN not necessary (105). The split is exploratory and outside the registered plan; no multiplicity correction.

	Cost			Forced inquiry			Targeted ask		
Model	N	P	NN	N	P	NN	N	P	NN
Main mean	+0.289	+0.149	+0.061	+0.035	+0.023	+0.007	+0.058	+0.015	+0.016
Main mean, artifacts excluded	+0.282	+0.158	+0.062	+0.034	+0.026	+0.006	+0.059	+0.017	+0.014
Qwen3-8B	+0.253	+0.124	+0.105	+0.034	+0.028	+0.000	+0.039	−0.004	+0.019
Qwen3-14B	+0.253	+0.183	+0.086	+0.050	+0.044	−0.010	+0.034	+0.040	+0.048
Meerkat2-14B	+0.292	+0.147	+0.067	−0.034	+0.000	−0.010	+0.050	+0.012	+0.029
MedGemma-4B	+0.168	+0.088	+0.048	+0.013	+0.008	−0.010	+0.018	−0.020	−0.029
Qwen3-30B-A3B	+0.300	+0.163	+0.048	+0.021	+0.000	−0.010	+0.058	+0.024	−0.010
Meerkat2-30B-A3B	+0.337	+0.120	+0.019	+0.037	+0.024	−0.038	+0.050	−0.004	+0.029
Qwen3-32B	+0.332	+0.183	+0.076	+0.082	+0.020	+0.010	+0.089	+0.052	+0.019
Meerkat2-32B	+0.345	+0.159	+0.029	−0.003	+0.032	−0.010	+0.105	+0.040	+0.019
MedGemma-27B	+0.326	+0.171	+0.076	+0.113	+0.056	+0.143	+0.082	+0.000	+0.019
Meerkat-70B	+0.308	+0.167	+0.048	+0.082	+0.064	+0.010	+0.082	+0.040	−0.019
DeepSeek-V4-Flash	+0.368	+0.171	+0.019	+0.111	+0.060	+0.010	+0.187	+0.056	−0.038

The larger cost on necessary items is partly built into the labels, which judge whether the removed facts decide the answer. 95% CIs for every cell are in results/p2_strata.json.

Table S9: Share of the accuracy lost by removal that each policy recovers, by necessity stratum (LLM labels; 380 necessary, 251 partially necessary, 105 not necessary items): sum over the stratum of (policy S1 − C0 S1) divided by the sum of (C0 S0 − C0 S1); main mean = ratio of the model-mean item vectors. These estimates are imprecise, and differences between strata were not formally tested. Post hoc, exploratory analysis; no multiplicity correction.

		Recovered share of the loss (%)			(S1−S2) / loss (%)	
Rows	Stratum	Free agent	Forced inquiry	Targeted ask	Forced inquiry	Targeted ask
Main mean (9 models)	necessary	1 [-3, 4]	12 [7, 17]	20 [16, 25]	14 [11, 18]	23 [19, 28]
	partially	2 [-7, 10]	16 [4, 27]	10 [1, 19]	13 [8, 19]	18 [11, 25]
	not necessary		not reported: small cost, wide interval			
Main mean, artifacts excluded	necessary	1 [-3, 5]	12 [7, 17]	21 [16, 26]	14 [11, 18]	22 [18, 27]
	partially	1 [-6, 9]	16 [5, 27]	11 [2, 19]	12 [8, 17]	17 [10, 24]
	not necessary		not reported: small cost, wide interval			
Meerkat-70B	necessary	-2 [-7, 3]	26 [13, 40]	26 [17, 36]	36 [25, 48]	27 [18, 37]
	partially	-7 [-24, 6]	38 [15, 65]	24 [7, 42]	36 [16, 62]	19 [8, 33]
	not necessary		not reported: small cost, wide interval			
DeepSeek-V4-Flash	necessary	6 [-3, 14]	30 [17, 42]	51 [40, 61]	5 [2, 9]	48 [38, 60]
	partially	19 [-3, 39]	35 [0, 65]	33 [9, 52]	0 [0, 0]	42 [23, 67]
	not necessary		not reported: small cost, wide interval			

Added after seeing the stratum differences in the cost of removal. 95% item-clustered bootstrap. DeepSeek-V4-Flash, partially necessary, forced inquiry: in all 251 items the model's outputs and final answers are identical at S1 and S2 (the simulator replies differ in 4 items, where a question was answered at S1, without changing any output), so (S1−S2)/loss is exactly 0 in every resample.

Table S10: Questions written by a Claude model that did not see the answer, the removed facts or any model output (blind asker, three questions per item), against the targeted ask pipeline of each evaluated model, on the same 200 items (%). Post hoc, exploratory analysis; no multiplicity correction.

Questions from	disclosed/question	≥1 fact	both facts	≥1 fact, necessary
Blind asker	42.5 [39.3, 45.7]	89.5 [85.0, 93.5]	40.5 [33.5, 47.5]	91.3
Qwen3-8B	17.3 [14.5, 20.3]	44.0 [37.0, 51.0]	13.5 [9.0, 18.5]	45.6
Qwen3-14B	22.8 [19.3, 26.5]	52.0 [45.0, 59.0]	15.0 [10.0, 20.0]	55.3
Meerkat2-14B	26.0 [22.5, 29.7]	56.5 [50.0, 63.5]	16.0 [11.0, 21.0]	59.2
MedGemma-4B	11.9 [9.1, 14.8]	28.5 [22.5, 35.0]	5.5 [2.5, 9.0]	32.0
Qwen3-30B-A3B	22.9 [19.6, 26.5]	54.0 [47.0, 61.0]	13.0 [8.5, 18.0]	60.2
Meerkat2-30B-A3B	27.1 [22.8, 31.6]	48.5 [41.5, 55.5]	18.5 [13.0, 24.0]	49.5
Qwen3-32B	27.8 [24.3, 31.3]	63.5 [57.0, 70.0]	18.5 [13.5, 24.0]	67.0
MedGemma-27B	22.5 [19.2, 26.0]	52.0 [45.5, 59.0]	16.5 [11.5, 21.5]	53.4
Meerkat2-32B	34.4 [30.2, 38.7]	62.5 [56.0, 69.0]	19.5 [14.0, 25.0]	64.1
Main mean, targeted ask	23.6 [21.3, 26.1]	51.3 [47.1, 55.6]	15.1 [12.3, 18.2]	54.0
Meerkat-70B	25.6 [22.2, 29.1]	55.0 [48.0, 62.0]	16.0 [11.0, 21.0]	53.4
DeepSeek-V4-Flash	29.7 [26.3, 33.0]	70.0 [63.5, 76.5]	26.5 [20.5, 32.5]	67.0

New exploratory comparison: the comparator is the targeted ask pipeline (three questions generated at once with the targeting instruction), not the ask pipeline of the earlier registered comparison. The subset oversamples necessary items (103 necessary, 68 partially necessary, 29 not necessary), so absolute rates are not population values. The blind asker is from the same Claude family as the item author and labeler, so it is a reference point rather than an independent standard. Its questions were sent to the same simulator weights in state S1, one at a time, on a server instance run before the evaluated models (replies not re-queried). 95% item-clustered bootstrap over the 200 items; main mean = mean of per-model values.

Table S11: Spearman rank correlation across models of per-model means, with 95% intervals from the shared item bootstrap. Post hoc, exploratory analysis, added after seeing the policy and acquisition tables; the unit is the model, so n is small, and model strength may raise both quantities. No multiplicity correction.

Pair	All ($n = 11$)	Main pool ($n = 9$)
Targeted ask: ≥1 fact vs Δ at S1	+0.85 [+0.56, +0.95]	+0.87 [+0.48, +0.96]
Targeted ask: ≥1 fact vs S1−S2	+0.85 [+0.50, +0.93]	+0.85 [+0.38, +0.94]
Forced inquiry: ≥1 fact vs S1−S2	+0.94 [+0.79, +0.98]	+0.95 [+0.72, +0.98]
Forced inquiry: ≥1 fact vs Δ at S1	+0.54 [+0.28, +0.68]	+0.94 [+0.52, +0.96]

For forced inquiry, S1−S2 can only arise when a fact was received, so its correlation with receipt is partly mechanical. DeepSeek-V4-Flash receives almost no facts in forced inquiry but gains at both S1 and S2, which lowers the correlation of receipt with Δ when the large models are included.

Table S12: MedGapQA construction and necessity labels. Each MedGapQA item is a MedQA case with two patient facts removed.

Group	Item	Count
MedQA test items	total	1,273
	excluded: no single patient	66
	excluded: image reference	114
	eligible	1,093
	authored as two-fact removals	750
	declined by the author	343
	excluded by leakage rule	14
MedGapQA items	total	736
	necessary	380
	partially	251
	not necessary	105
	broken	0
	artifact flag	44

Table S13: Type of each removed fact and whether a patient could report it (two facts per item). Post hoc analysis (not pre-specified); no multiplicity correction.

		Patient could report		
Fact type	Facts	Yes	Partly	No
Symptom or history	483	465	18	0
Examination finding	394	1	124	269
Laboratory value	231	0	8	223
Imaging or procedure	138	0	3	135
Medication	70	53	17	0
Other	70	16	7	47
Vital sign	47	0	9	38
Demographic	39	39	0	0
All facts	1472	574 (39.0%)	186 (12.6%)	712 (48.4%)

Labels from one Claude model, shown only the fact text, without clinician review; "patient could report" is a judgment, not a measurement.

Table S14: Audit of 120 simulator replies that withheld information ("not available" or "already given") at S1. Post hoc analysis (not pre-specified); no multiplicity correction.

Quantity	n	Unweighted (%)	Weighted (%) [95% CI]
Question asks for a patient fact	93	77.5	77.7 [67.7, 86.6]
Question is broad	26	21.7	18.0 [9.9, 27.1]
Targets a removed fact: yes	1	0.8	1.9 [0.0, 6.1]
Targets a removed fact: partly	9	7.5	6.3 [1.2, 12.9]
Correct reply: not available	86	71.7	74.5 [64.0, 83.8]
Correct reply: already given in the case	24	20.0	17.3 [9.7, 26.0]
Correct reply: ambiguous	9	7.5	6.3 [1.2, 12.9]
Correct reply: disclose a removed fact	1	0.8	1.9 [0.0, 6.1]
Simulator reply matches the correct reply	98	81.7	80.6 [70.4, 89.8]
Wrong refusal, lower bound (disclose)	1	0.8	1.9 [0.0, 6.1]
Wrong refusal, upper bound (disclose or ambiguous)	10	8.3	8.2 [2.1, 15.7]
Question clinically reasonable	91	75.8	74.8 [64.3, 84.4]
Over-disclosure (earlier audit): disclosure to an untargeted question	6	11.5	–

Rows sampled from 48 cells (four models $\times$ four asking policies $\times$ three necessity strata, two or three per cell); weights are inverse inclusion probabilities (questions in the cell's pool), CIs from 10,000 row resamples. 98 rows had the reply "not available" and 22 "already given". The over-disclosure row is from the earlier audit of 52 disclosing replies. Labels from one Claude model without clinician review.